

Assignment 5

1. Choose a REGRESSION dataset (reusing bikeshare is allowed), perform a test/train split, and build a regression model (just like in assignment 3), and calculate the

- + Training Error (MSE, MAE)
- + Testing Error (MSE, MAE)

```
In [111... import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
plt.rcParams['figure.figsize'] = 20, 10
import pandas as pd
import numpy as np
credit = pd.read_csv('../data/Credit.csv')
credit.head()
```

```
Out[111... Unnamed:
           0  Income  Limit  Rating  Cards  Age  Education  Gender  Student  Married
0          1  14.891  3606    283     2   34         11    Male     No      Yes
1          2  106.025 6645    483     3   82         15  Female     Yes     Yes
2          3  104.593 7075    514     4   71         11    Male     No      No
3          4  148.924 9504    681     3   36         11  Female     No      No
4          5   55.882 4897    357     2   68         16    Male     No      Yes
```

```
In [112... x = credit['Balance'].to_numpy()
y = credit['Rating'].to_numpy()

x = x.reshape(-1, 1)
y = y.reshape(-1, 1)
```

```
In [113... from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=.40)
```

```
In [114... from sklearn import linear_model, metrics
from sklearn.linear_model import LinearRegression
model = LinearRegression()
```


```
In [115... model.fit(x_train, y_train)
model.coef_, model.intercept_
```

Out[115...] (array([[0.29742749]]), array([202.35136259]))

Plot of Training Model (Red) versus Actual Data (Blue)

```
In [116...] plt.scatter(x, y)
plt.plot(x, np.dot(x, model.coef_) + model.intercept_, c='r')
```

Out[116...] [<matplotlib.lines.Line2D at 0x1ba6939e290>]



Training Error (MSE, MAE)

```
In [117...] (
    metrics.mean_squared_error(y_train, np.dot(x_train, model.coef_) + model.intercept_),
    metrics.mean_absolute_error(y_train, np.dot(x_train, model.coef_) + model.intercept_)
)
```

Out[117...] (6435.394001360988, 59.73485084756984)

Testing Error (MSE, MAE)

```
In [118...] (
    metrics.mean_squared_error(y_test, np.dot(x_test, model.coef_) + model.intercept_),
    metrics.mean_absolute_error(y_test, np.dot(x_test, model.coef_) + model.intercept_)
)
```

Out[118...] (5554.666371103305, 55.07999097492727)

2. Choose a CLASSIFICATION dataset (not the adult.data set, The UCI repository has many datasets as well as Kaggle), perform test/train split

and create a classification model (your choice but DecisionTree is fine). Calculate

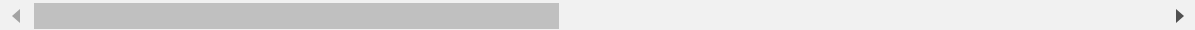
- + Accuracy
- + Confusion Matrix
- + Classification Report

```
In [119...] df = pd.read_csv('../data/College.csv', index_col=False)
```

```
In [120...] df.head()
```

```
Out[120...]      Unnamed: 0  Private  Apps  Accept  Enroll  Top10perc  Top25perc  F.Undergrad  P.Undergrad
```

0	Abilene Christian University	Yes	1660	1232	721	23	52	2885	
1	Adelphi University	Yes	2186	1924	512	16	29	2683	1
2	Adrian College	Yes	1428	1097	336	22	50	1036	
3	Agnes Scott College	Yes	417	349	137	60	89	510	
4	Alaska Pacific University	Yes	193	146	55	16	44	249	



```
In [121...] df['Private'].unique()
```

```
Out[121...] array(['Yes', 'No'], dtype=object)
```

```
In [122...] df.columns
```

```
Out[122...] Index(['Unnamed: 0', 'Private', 'Apps', 'Accept', 'Enroll', 'Top10perc',  
                'Top25perc', 'F.Undergrad', 'P.Undergrad', 'Outstate', 'Room.Board',  
                'Books', 'Personal', 'PhD', 'Terminal', 'S.F.Ratio', 'perc.alumni',  
                'Expend', 'Grad.Rate'],  
                dtype='object')
```

```
In [123...] non_numeric_columns = ['Unnamed: 0']  
x1 = df.copy().drop(non_numeric_columns, axis = 1)  
x1['Private'] = x1.Private.str.contains('Yes').astype(int)  
x1.head()
```

Out[123...

	Private	Apps	Accept	Enroll	Top10perc	Top25perc	F.Undergrad	P.Undergrad	Outsta
0	1	1660	1232	721	23	52	2885	537	74
1	1	2186	1924	512	16	29	2683	1227	122
2	1	1428	1097	336	22	50	1036	99	112
3	1	417	349	137	60	89	510	63	129
4	1	193	146	55	16	44	249	869	75

In [124... `x1.Private.value_counts()`

Out[124... Private
1 565
0 212
Name: count, dtype: int64

In [125... `x1_train, x1_test, y1_train, y1_test = train_test_split(x1.drop(['Private'], axis=1`

In [126... `x1.shape, x1_train.shape, x1_test.shape`

Out[126... `((777, 18), (621, 17), (156, 17))`

In [127... `from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import (accuracy_score,
classification_report,
confusion_matrix, auc, roc_curve
)
model1 = DecisionTreeClassifier(criterion='entropy')`

In [128... `model1.fit(x1_train, y1_train)`

Out[128... `DecisionTreeClassifier`
`DecisionTreeClassifier(criterion='entropy')`

In [129... `test_predictions = model1.predict(x1_test)`

Accuracy Score

In [130... `accuracy_score(y1_test, test_predictions)`

Out[130... `0.9230769230769231`

Confusion Matrix

In [131... `confusion_matrix(y1_test, test_predictions)`

```
Out[131...] array([[ 34,   3],
        [  9, 110]], dtype=int64)
```

Classification Report

```
In [132...] print(classification_report(y1_test, test_predictions))
```

	precision	recall	f1-score	support
0	0.79	0.92	0.85	37
1	0.97	0.92	0.95	119
accuracy			0.92	156
macro avg	0.88	0.92	0.90	156
weighted avg	0.93	0.92	0.92	156

3. (Bonus) See if you can improve the classification model's performance with any tricks you can think of (modify features, remove features, polynomial features)

```
In [133...] list(zip(x1.drop(['Private'], axis=1).columns, model1.feature_importances_))
```

```
Out[133...] [('Apps', 0.007508158360841993),
 ('Accept', 0.0073207069197507735),
 ('Enroll', 0.055996120071413746),
 ('Top10perc', 0.04299202955266883),
 ('Top25perc', 0.0),
 ('F.Undergrad', 0.4019146762155514),
 ('P.Undergrad', 0.04157049674065695),
 ('Outstate', 0.30405059366229037),
 ('Room.Board', 0.018979755625197606),
 ('Books', 0.0),
 ('Personal', 0.017353442174389035),
 ('PhD', 0.039485948188319145),
 ('Terminal', 0.027007203544034727),
 ('S.F.Ratio', 0.012182409266252097),
 ('perc.alumni', 0.015264108687949503),
 ('Expend', 0.00837435099068376),
 ('Grad.Rate', 0.0)]
```

```
In [134...] x1.columns
```

```
Out[134...] Index(['Private', 'Apps', 'Accept', 'Enroll', 'Top10perc', 'Top25perc',
                  'F.Undergrad', 'P.Undergrad', 'Outstate', 'Room.Board', 'Books',
                  'Personal', 'PhD', 'Terminal', 'S.F.Ratio', 'perc.alumni', 'Expend',
                  'Grad.Rate'],
                  dtype='object')
```

```
In [135...] #Low importance features
drop_columns = ['Apps', 'Accept', 'Enroll', 'Top25perc', 'Books',
                'Personal', 'Terminal', 'S.F.Ratio', 'Expend']
```

```
x2 = x1.copy().drop(drop_columns, axis = 1)
x2.head()
```

```
Out[135...]      Private  Top10perc  F.Undergrad  P.Undergrad  Outstate  Room.Board  PhD  perc.alumni
0          1         23         2885         537        7440         3300   70         12
1          1         16         2683        1227       12280         6450   29         16
2          1         22         1036          99       11250         3750   53         30
3          1         60          510          63       12960         5450   92         37
4          1         16          249         869       7560         4120   76          2
```



```
In [136...] x2_train, x2_test, y2_train, y2_test = train_test_split(x2.drop(['Private'], axis=1),
```

```
In [137...] model2 = DecisionTreeClassifier(criterion='entropy')
model2.fit(x2_train, y2_train)
```

```
Out[137...] DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy')
```

```
In [138...] test2_predictions = model2.predict(x2_test)
```

Accuracy

```
In [139...] accuracy_score(y2_test, test2_predictions)
```

```
Out[139...] 0.9487179487179487
```

Confusion Matrix

```
In [140...] confusion_matrix(y2_test, test2_predictions)
```

```
Out[140...] array([[ 41,   3],
       [  5, 107]], dtype=int64)
```

Classification Report

```
In [141...] print("Test 2", classification_report(y2_test, test2_predictions), "Test 1", classi
```

Test 2		precision	recall	f1-score	support
	0	0.89	0.93	0.91	44
	1	0.97	0.96	0.96	112
	accuracy			0.95	156
	macro avg	0.93	0.94	0.94	156
	weighted avg	0.95	0.95	0.95	156
Test 1		precision	recall	f1-score	support
	0	0.79	0.92	0.85	37
	1	0.97	0.92	0.95	119
	accuracy			0.92	156
	macro avg	0.88	0.92	0.90	156
	weighted avg	0.93	0.92	0.92	156