

# Neural Networks image recognition - MultiLayer Perceptron

Use both MLNN for the following problem.

1. Add random noise (see below on `size` parameter on `np.random.normal`) to the images in training and testing. **\*\*Make sure each image gets a different noise feature added to it. Inspect by printing out several images. Note - the `size` parameter should match the data. \*\***
2. Compare the `accuracy` of train and val after N epochs for MLNN with and without noise.

```
In [1]: #setup
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.optimizers import RMSprop
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: # data import
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# model setup
batch_size = 128
num_classes = 10
epochs = 20
```

60000 train samples

10000 test samples

```
In [3]: #build model

model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer="adam",
              metrics=['accuracy'])
```

```
C:\Users\rober\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 512)	401,920
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 512)	262,656
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 10)	5,130

Total params: 669,706 (2.55 MB)

Trainable params: 669,706 (2.55 MB)

Non-trainable params: 0 (0.00 B)

```
In [4]: #add noise
np.random.seed(15)
noise_feat = np.linspace(.1, 1.0, 10)
np.random.shuffle(noise_feat)
feat_img = {i: noise_feat[i] for i in range(10)}
x_train_noise = []
x_test_noise = []
#noise for train
for i in range(len(x_train)):
    class_lab = y_train[i]
    noise_scale = feat_img[class_lab]
    noise = np.random.normal(loc=0.0, scale=noise_scale, size=x_train[i].shape)
    img = np.clip(x_train[i] + noise, 0., 1.)
    x_train_noise.append(img)
#noise for test
for i in range(len(x_test)):
    class_lab = y_test[i]
    noise_scale = feat_img[class_lab]
    noise = np.random.normal(loc=0.0, scale=noise_scale, size=x_test[i].shape)
    img = np.clip(x_test[i] + noise, 0., 1.)
    x_test_noise.append(img)


x_train_noise = np.array(x_train_noise)
x_test_noise = np.array(x_test_noise)
```

```
In [5]: #compare plots with noise
plt.figure()
n = 10
for i in range(n):
    ax = plt.subplot(2, n, i+1)
    plt.imshow(x_train[i].reshape(28, 28))
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    ax = plt.subplot(2, n, i+n+1)
    plt.imshow(x_train_noise[i].reshape(28, 28))
    ax.get_xaxis().set_visible(False)
```

```
ax.get_yaxis().set_visible(False)
```

```
plt.show()
```



### Train Data (no noise)

```
In [6]: # convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
In [7]: history_orig = model.fit(x_train, y_train,
                                batch_size=batch_size,
                                epochs=epochs,
                                verbose=1,
                                validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Epoch 1/20  
469/469 ————— 3s 4ms/step - accuracy: 0.8648 - loss: 0.4566 - val\_accuracy: 0.9671  
- val\_loss: 0.1047  
Epoch 2/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9659 - loss: 0.1071 - val\_accuracy: 0.9754  
- val\_loss: 0.0850  
Epoch 3/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9774 - loss: 0.0713 - val\_accuracy: 0.9785  
- val\_loss: 0.0701  
Epoch 4/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9823 - loss: 0.0543 - val\_accuracy: 0.9783  
- val\_loss: 0.0702  
Epoch 5/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9848 - loss: 0.0454 - val\_accuracy: 0.9796  
- val\_loss: 0.0724  
Epoch 6/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9887 - loss: 0.0360 - val\_accuracy: 0.9800  
- val\_loss: 0.0695  
Epoch 7/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9898 - loss: 0.0304 - val\_accuracy: 0.9811  
- val\_loss: 0.0657  
Epoch 8/20  
469/469 ————— 3s 6ms/step - accuracy: 0.9893 - loss: 0.0305 - val\_accuracy: 0.9802  
- val\_loss: 0.0766  
Epoch 9/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9916 - loss: 0.0252 - val\_accuracy: 0.9799  
- val\_loss: 0.0819  
Epoch 10/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9913 - loss: 0.0249 - val\_accuracy: 0.9817  
- val\_loss: 0.0697  
Epoch 11/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9929 - loss: 0.0220 - val\_accuracy: 0.9811  
- val\_loss: 0.0794  
Epoch 12/20  
469/469 ————— 3s 6ms/step - accuracy: 0.9927 - loss: 0.0224 - val\_accuracy: 0.9822  
- val\_loss: 0.0695  
Epoch 13/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9940 - loss: 0.0175 - val\_accuracy: 0.9798  
- val\_loss: 0.0842  
Epoch 14/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9944 - loss: 0.0169 - val\_accuracy: 0.9815  
- val\_loss: 0.0731  
Epoch 15/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9941 - loss: 0.0170 - val\_accuracy: 0.9816  
- val\_loss: 0.0721  
Epoch 16/20  
469/469 ————— 3s 6ms/step - accuracy: 0.9945 - loss: 0.0171 - val\_accuracy: 0.9826  
- val\_loss: 0.0716  
Epoch 17/20  
469/469 ————— 3s 5ms/step - accuracy: 0.9952 - loss: 0.0138 - val\_accuracy: 0.9794  
- val\_loss: 0.0889  
Epoch 18/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9956 - loss: 0.0135 - val\_accuracy: 0.9824  
- val\_loss: 0.0879  
Epoch 19/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9952 - loss: 0.0154 - val\_accuracy: 0.9829  
- val\_loss: 0.0793  
Epoch 20/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9950 - loss: 0.0153 - val\_accuracy: 0.9836  
- val\_loss: 0.0826  
Test loss: 0.08255087584257126  
Test accuracy: 0.9836000204086304

## Train Data (w/ noise)

```
In [8]: history2 = model.fit(x_train_noise, y_train,
                             batch_size=batch_size,
                             epochs=epochs,
                             verbose=1,
                             validation_data=(x_test_noise, y_test))
score2 = model.evaluate(x_test_noise, y_test, verbose=0)
print('Test loss:', score2[0])
print('Test accuracy:', score2[1])
```

Epoch 1/20  
469/469 ————— 2s 4ms/step - accuracy: 0.8310 - loss: 0.5252 - val\_accuracy: 0.9288  
- val\_loss: 0.2050  
Epoch 2/20  
469/469 ————— 2s 4ms/step - accuracy: 0.9334 - loss: 0.1852 - val\_accuracy: 0.9354  
- val\_loss: 0.1841  
Epoch 3/20  
469/469 ————— 3s 6ms/step - accuracy: 0.9501 - loss: 0.1333 - val\_accuracy: 0.9420  
- val\_loss: 0.1743  
Epoch 4/20  
469/469 ————— 3s 6ms/step - accuracy: 0.9642 - loss: 0.1014 - val\_accuracy: 0.9410  
- val\_loss: 0.1756  
Epoch 5/20  
469/469 ————— 3s 6ms/step - accuracy: 0.9697 - loss: 0.0837 - val\_accuracy: 0.9439  
- val\_loss: 0.1738  
Epoch 6/20  
469/469 ————— 3s 6ms/step - accuracy: 0.9762 - loss: 0.0662 - val\_accuracy: 0.9399  
- val\_loss: 0.1913  
Epoch 7/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9799 - loss: 0.0556 - val\_accuracy: 0.9379  
- val\_loss: 0.2074  
Epoch 8/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9842 - loss: 0.0462 - val\_accuracy: 0.9414  
- val\_loss: 0.2174  
Epoch 9/20  
469/469 ————— 3s 5ms/step - accuracy: 0.9866 - loss: 0.0374 - val\_accuracy: 0.9379  
- val\_loss: 0.2424  
Epoch 10/20  
469/469 ————— 3s 6ms/step - accuracy: 0.9856 - loss: 0.0410 - val\_accuracy: 0.9346  
- val\_loss: 0.2706  
Epoch 11/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9886 - loss: 0.0350 - val\_accuracy: 0.9354  
- val\_loss: 0.2695  
Epoch 12/20  
469/469 ————— 3s 6ms/step - accuracy: 0.9891 - loss: 0.0323 - val\_accuracy: 0.9408  
- val\_loss: 0.2433  
Epoch 13/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9896 - loss: 0.0297 - val\_accuracy: 0.9370  
- val\_loss: 0.2901  
Epoch 14/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9897 - loss: 0.0309 - val\_accuracy: 0.9380  
- val\_loss: 0.2643  
Epoch 15/20  
469/469 ————— 3s 5ms/step - accuracy: 0.9906 - loss: 0.0296 - val\_accuracy: 0.9380  
- val\_loss: 0.2561  
Epoch 16/20  
469/469 ————— 3s 6ms/step - accuracy: 0.9912 - loss: 0.0262 - val\_accuracy: 0.9431  
- val\_loss: 0.2710  
Epoch 17/20  
469/469 ————— 3s 6ms/step - accuracy: 0.9920 - loss: 0.0238 - val\_accuracy: 0.9346  
- val\_loss: 0.2996  
Epoch 18/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9903 - loss: 0.0288 - val\_accuracy: 0.9380  
- val\_loss: 0.3006  
Epoch 19/20  
469/469 ————— 2s 5ms/step - accuracy: 0.9924 - loss: 0.0221 - val\_accuracy: 0.9390  
- val\_loss: 0.3031  
Epoch 20/20  
469/469 ————— 3s 6ms/step - accuracy: 0.9910 - loss: 0.0275 - val\_accuracy: 0.9396  
- val\_loss: 0.2976  
Test loss: 0.2975553274154663  
Test accuracy: 0.9395999908447266

## Comparison

```
In [9]: print('Test loss:', score[0], 'Test loss with Noise:', score2[0])
        print('Test accuracy:', score[1], 'Test accuracy with Noise:', score2[1])
```

Test loss: 0.08255087584257126 Test loss with Noise: 0.2975553274154663

Test accuracy: 0.9836000204086304 Test accuracy with Noise: 0.9395999908447266

## 3. Vary the amount of noise by changing the scale parameter in np.random.normal by a factor. Use .1, .5, 1.0, 2.0, 4.0 for the scale and keep track of the accuracy for training and validation and plot these results

```
In [10]: ##Build Models
```

```
def noise_test(df, scale):
    noise = np.random.normal(loc=0.0, scale=scale, size=df.shape)
    df_noise = df + noise
    return np.clip(df_noise, 0., 1.)

def mlnn():
    model = Sequential()
    model.add(Dense(512, activation='relu', input_shape=(784,)))
    model.add(Dropout(0.2))
    model.add(Dense(512, activation='relu'))
    model.add(Dropout(0.2))
    model.add(Dense(10, activation='softmax'))
    model.compile(loss='categorical_crossentropy',
                  optimizer="adam",
                  metrics=['accuracy'])
    return model

def eval(x_train, y_train, x_test, y_test, scale):
    x_train_noise = noise_test(x_train, scale)
    x_test_noise = noise_test(x_test, scale)

    model = mlnn()
    history = model.fit(x_train_noise, y_train,
                       epochs = 20,
                       batch_size = 128,
                       validation_data=(x_test_noise, y_test),
                       verbose = 0)

    return history

#Citation: OpenAI (2024). ChatGPT. Retrieved from: openai.chatgpt.com
```

```
In [11]: #test
noise_scale = [0.1, 0.5, 1.0, 2.0, 4.0]
histories = {}

for scale in noise_scale:
    print(f'Training Noise Scale: {scale}')
    histories[scale] = eval(x_train, y_train, x_test, y_test, scale)
```

Training Noise Scale: 0.1  
Training Noise Scale: 0.5  
Training Noise Scale: 1.0  
Training Noise Scale: 2.0  
Training Noise Scale: 4.0

```
In [12]: #plot for training and validation
fig, axs = plt.subplots(1, 2)


for scale, history in histories.items():
    axs[0].plot(history.history['accuracy'], label=f'Noise Scale={scale}')

axs[0].set_title('Training Accuracy')
axs[0].set_xlabel('Epoch')
axs[0].set_ylabel('Accuracy')

for scale, history in histories.items():
    axs[1].plot(history.history['val_accuracy'])

axs[1].set_title('Validation Accuracy')
axs[1].set_xlabel('Epoch')
axs[1].set_ylabel('Accuracy')

axs[0].set_ylim(0.1, 1.0)
axs[1].set_ylim(0.1, 1.0)
fig.legend(loc='center left', bbox_to_anchor=(1, 0.5))
plt.show()
```



In [ ]: